

CHAPTER 10

Mixing C and Assembly

Occasionally it is required to write a program in both C and assembly language. There are several possible reasons.

- First, an experienced programmer might want to optimize a performance-critical function manually in assembly, instead of relying on compilers. Many profiling tools can identify the most time-consuming functions and compilers often have limited intelligence in optimizing these functions. A handcrafted assembly code can out-perform high-level languages, such as C.
- Second, writing a program in assembly allows a programmer to use processor-specific instructions. For example, a test-and-set atomic assembly instruction can implement locks and semaphores. Another example is that most C compilers do not use some operations available on Cortex-M processors, such as ROR (rotate right) and RRX (rotate right extended).
- Third, assembly programs can directly access hardware, which is especially useful for device drivers and processor booting code.

The embedded application binary interface (EABI) briefly introduced in Chapter 8.2 defines low-level standards of interfacing program modules that are compiled separately, no matter whether these modules are written in C or assembly. The EABI specifies (1) standards for data types, data alignments, and executable file formats, and (2) conventions for function calls, parameter passing, registers usage, and stack frame. If a program is written in C, compilers ensure that these standards are followed strictly. However, if a program is developed in assembly, it is

"The good thing about standards is that there are so many to choose from."

Andrew Tanenbaum,
famous computer
scientist

the programmer's responsibility to adhere to these standards. The standard allows programmers to mix C and assembly in the application.

10.1 Data Types and Access

While the size of a basic data type in the C language depends on the compilers and platforms, the following table lists the typical size of commonly used data types in the C language.

Data Type	Size (bits)	Alignment	Data Range
bool	8	byte	0 or 1. Bits 1 – 7 are ignored
char	8	byte	-128 to 127 (signed), or 0 to 255 (unsigned)
int	32	word	-2,147,483,648 to 2,147,483,647 (signed), or 0 to 4,294,967,296 (unsigned)
short int	16	halfword	-32,768 to 32,767 (signed), or 0 to 65,536 (unsigned)
long int	32	word	same as int
long long	64	word	-9,223,372,036,854,775,808 to 9,223,372,036,854,775,807 (signed), or 0 to 18,446,744,073,709,551,616 (unsigned)
float	32	word	+/- (1.4023×10^{-45} to 3.4028×10^{38}), always signed
double	64	word	+/- (4.9406×10^{-324} to 1.7977×10^{308}), always signed
long double	96	word	very large range
pointer	32	word	0 to 4,294,967,296

Table 10-1. Data size and alignment of basic data types in C

10.1.1 Signed or Unsigned Integers

When programming in an assembly language, it is the programmer's responsibility to interpret whether a data item is signed or unsigned. For example, when loading an 8-bit data into a 32-bit register, the program should use LDRSB (load register with signed byte) to access a signed character and LDRB (load register with byte) to retrieve an unsigned character.

- LDRSB loads a byte from the memory into a register and performs sign extension. The sign extension duplicates the sign bit of the 8-bit data to all bits at the most significant side of a register to preserve the positive or negative sign.
- LDRB loads a byte from the memory into a register and simply pads the left of the register with zeros.

For example, when a program loads an 8-bit binary data 0x88 (+136 for unsigned or -120 for signed) from the memory into a 32-bit register, should the register be 0xFFFFFFF88 or

`0x00000088`? It depends on the programmer's intention. If these 8 bits represent a signed number, `LDRSB` should be used to preserve the sign. If they represent an unsigned number, `LDRB` should be used. Similarly, `LDRSH` (load register with signed halfword) and `LDRH` (load register with halfword) bring a 16-bit signed and unsigned number into a register, respectively. Table 10-2 summarizes these load instructions.

Variable	Instruction	Description	Sign Extension
<code>unsigned char</code>	<code>LDRB</code>	Load register with byte	No
<code>unsigned short</code>	<code>LDRH</code>	Load register with halfword	No
<code>unsigned/signed int</code>	<code>LDR</code>	Load register with word	No
<code>char</code>	<code>LDRSB</code>	Load register with signed byte	Yes
<code>short</code>	<code>LDRSH</code>	Load register with signed halfword	Yes

Table 10-2. ARM assembly instructions for accessing various basic integer data types

Correspondingly, `STRB` (store register byte) and `STRH` (store register halfword) store either a signed number or an unsigned number into the memory. Loading or storing a 32-bit integer does not need to take care of the sign because each register has the same amount of bits as the integer.

A 64-bit integer takes two registers, and it can be loaded by using two separate `LDR` instructions or a single `LDRD` (load registers with double words).

C Program	Assembly Program
<code>signed long long x = -1;</code>	<pre> LDR r3, =x ; load memory address of x LDRD r0, r1, [r3] ; r0 lower word, r1 higher word x DCW 0xFFFFFFFF, 0xFFFFFFFF ; allocate 8 bytes </pre>

Example 10-1. Loading a 64-bit integer from the memory

10.1.2 Data Alignment

Most computer systems have some alignment requirement on the starting memory address of a variable. The memory address of a C variable often has to be aligned, as listed in Table 10-1. The smallest unit exchanged between the processor and the memory is a byte (8 bits), and thus the memory address is always in terms of bytes.

A variable is n -byte aligned in memory if its starting memory address is some multiple of n . Typically, n is a power of 2, such as 2 (halfword aligned), 4 (word aligned), and 8 (double word aligned). Suppose a 32-bit variable is word aligned. If the address of the next available byte in memory is `0x8001`, the variable is then stored in a continuous span of 4 bytes from `0x8004` to `0x8007`. The compiler or the program inserts three meaningless

bytes at memory addresses $0x8001$, $0x8002$, and $0x8003$. These three bytes are called *padding bytes*.

Enforcing data alignment is to improve the memory performance. A memory system consists of multiple storage units, and the processor typically distributes data among these units in a round-robin fashion. Because the number of pins available on a processor is limited, these memory units have to share some pins in the memory address bus. To allow these memory units to transfer data concurrently, the target data stored in all memory units needs to share a portion of their memory addresses. The data alignment ensures that all data of a variable stored in different memory units meet this requirement. When the processor reads a properly aligned variable, only one access is required to transfer the data out of these memory units. Otherwise, two separate memory accesses might be necessary, slowing down the processor performance.

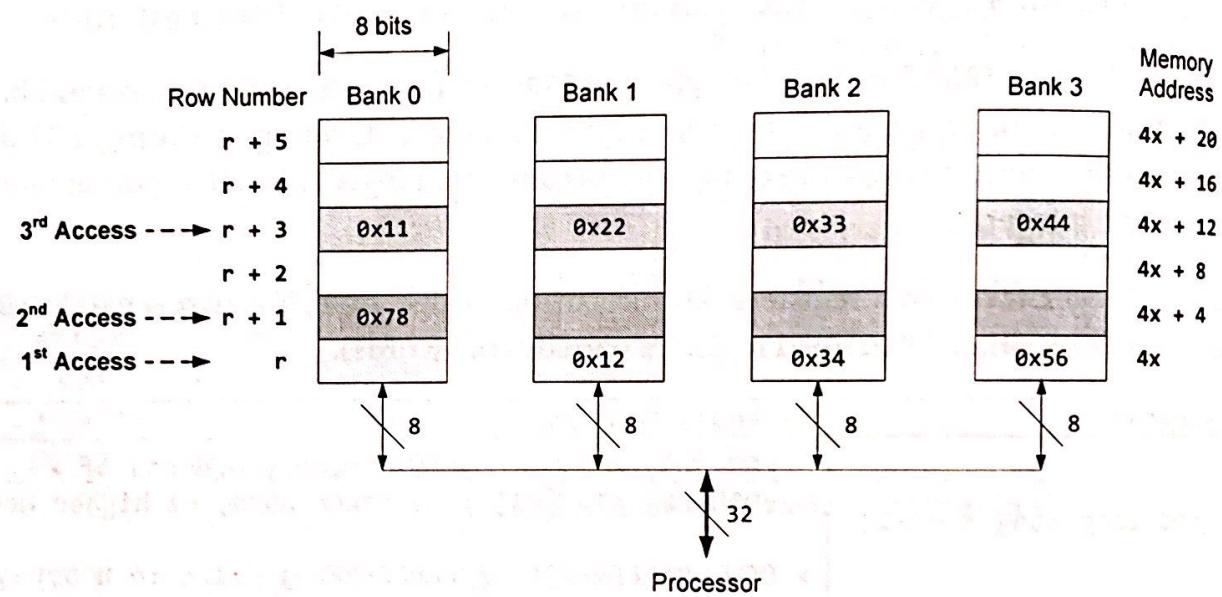


Figure 10-1. Loading unaligned data $0x78563412$ takes two accesses even if the processor allows unaligned memory accesses. Loading aligned data $0x44332211$ takes only one access.

As shown in Figure 10-1, the data memory is organized into four banks, and these banks can feed the 32-bit data bus. Four bytes in the same row of all banks can be loaded into the processor concurrently. In this example, data $0x78563412$ is not aligned with word boundaries, and the processor takes two memory accesses to load the data $0x78563412$ to a register. However, it takes only one memory access to load data $0x44332211$.

As introduced in Chapter 3.6, the “ALIGN” directive gives data alignment requirements to compilers. The syntax is “ALIGN boundary, offset”. The boundary is any power of 2. The default boundary is 4, making the next variable align to a word boundary. The offset specifies how many bytes the next variable should start from the word boundary. The default offset is 0.

```

AREA myData, DATA, ALIGN=2
; word aligned

a DCB 0x11

ALIGN 4
; word aligned
b DCD 0x12345678

ALIGN 4,3
; word aligned with an offset of 3
c DCB 0x22

ALIGN 4,2
; word aligned with an offset of 2
d DCW 0xAABB

```

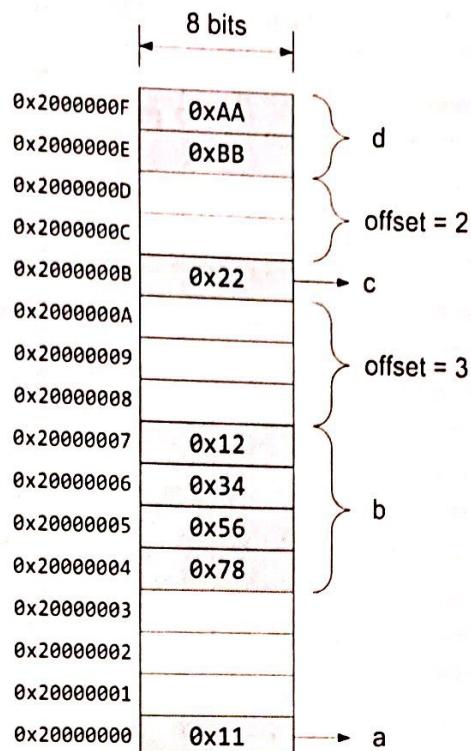


Figure 10-2. Memory layout

10.1.3 Data Structure Padding

A data structure defined in C language aggregates multiple basic variables into a single complex entity. By default, compilers ensure that all variables in a structure are aligned to their required memory boundaries. In a structure array, compilers also ensure that all variables in this array meet their alignment requirements. Therefore, compilers may place padding bytes between structure variables.

C language also supports *packed structures* in which variables are not aligned. Therefore, compilers do not add any padding bytes into a data structure. Packed structures are often used in communication protocols (such as USB) to save transmission time.

Figure 10-3 and Figure 10-4 compare the memory layout of an unpacked structure defined in Example 10-2.

Unpacked Structure	Packed Structure
<pre> struct Position { char x; char y; char z; int time; short scale; } array[2]; </pre>	<pre> __packed struct Position { char x; char y; char z; int time; short scale; } array[2]; </pre>

Example 10-2. Comparison of unpacked and packed structure in C

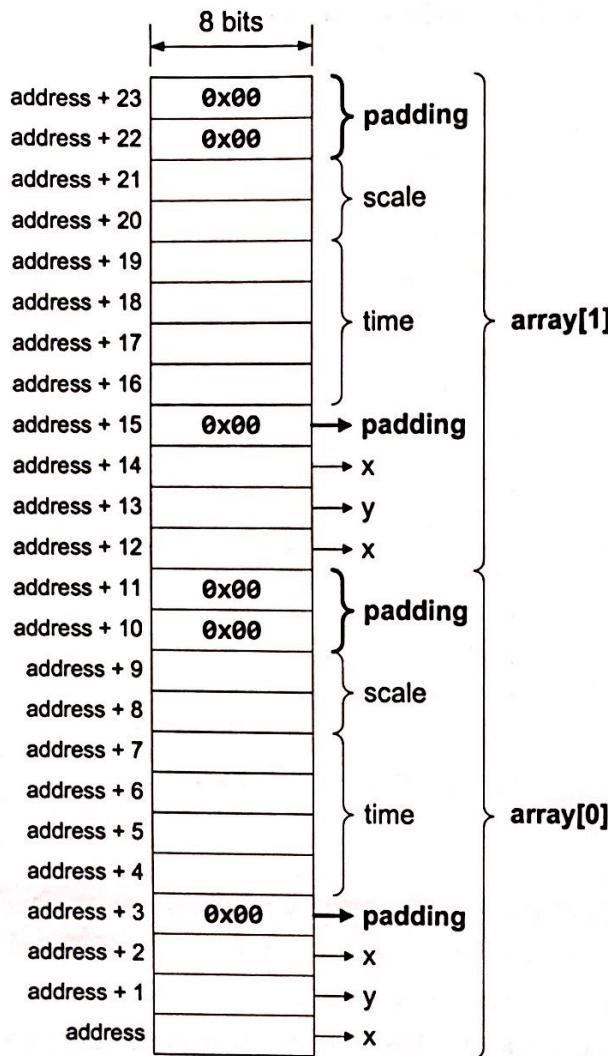


Figure 10-3. In an unpacked structure, variables are aligned. Specifically the integer variable and the structure are aligned in words.

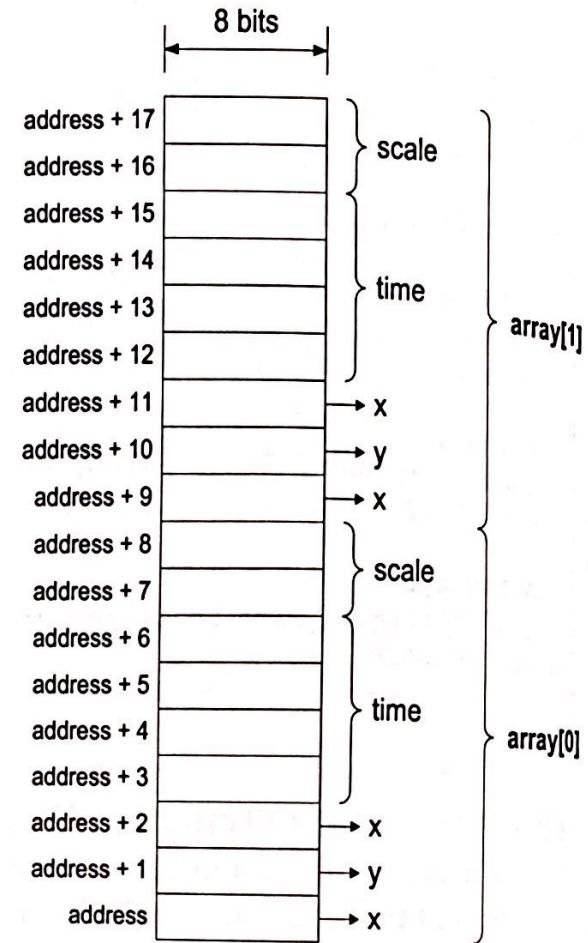


Figure 10-4. In a packed structure, variables are not aligned. Cortex-M processors support unaligned access in LDR/STR, LDRT/STRT, LDRH/STRH, and LDRHT/STRHT.

In Figure 10-3, the compiler inserts three padding bytes into the structure position.

- The first padding byte is inserted after variable *x* to make the next integer variable *time* aligned to some word boundary.
- In a structure array, the compiler also ensures that all variables in this array meet their alignment requirements. Therefore, two additional bytes are added at the end of the data structure, making the size of the *Position* structure a multiple of four. This padding also makes the variables in this array, particularly the *time* variable, align properly.

In fact, two structure definitions give below are equivalent. Compilers allocate proper padding bytes in an unpacked structure.

```
struct Position {
    char x;
    char y;
    char x;
    int time;
    short scale;
} array[2];
```

Equivalent
↔

```
struct Position {
    char x;
    char y;
    char x;
    char padding_1;
    int time;
    short scale;
    char padding_2[2];
} array[2];
```

Example 10-3. Compiler inserts three padding bytes to an unpacked structure

In Figure 10-4, the position structure uses the type modifier “`_packed`” to the compiler to produce an unaligned memory layout. Specifically, the integer variable `time` is not aligned to a word boundary, and the short variable `scale` is not aligned to a halfword boundary. Therefore, there is no padding between structure members or at the end of the structure. The `_unpack` modifier is often used to map a structure to a special data area in memory, such as a USB communication package received in a memory buffer.

Programs often do not use packed structures. No ARM processors released before ARM-V6 support unaligned memory accesses. The instruction “`LDR r1, [r0]`” would generate an alignment exception if the memory address stored in `r0` were not a multiple of four. The Cortex-M processors do support unaligned memory accesses. However, unaligned accesses are still slower than aligned memory accesses, and thus it is recommended to avoid using unaligned accesses.

Packed structures and unpacked structures are not compatible with each other. We cannot assign or cast one to the other. The only way to assign a packed structure to an unpacked structure is to copy all structure members one by one.

Suppose we want to set `array[0].time` to 1234. Example 10-4 compares the assembly codes that update the `time` variable of the unpacked and packed structure respectively.

Unpacked Structure	Packed Structure
<code>LDR r0, =array ; load base address</code>	<code>LDR r0, =array ; load base address</code>
<code>LDR r1, [r0, #4] ; array[0].time</code>	<code>LDR r1, [r0, #3] ; array[0].time</code>
<code>LDR r2, #1234</code>	<code>LDR r2, #1234</code>
<code>STR r2, [r0, #4] ; array[0].time</code>	<code>STR r2, [r0, #3] ; array[0].time</code>

Example 10-4. Accessing members of an unpacked and packed structure in assembly

In the unpacked structure, the access to `array[0].time` is aligned. However, in the packed structure, the access is misaligned. Even though the misaligned accesses “`LDR r1, [r0, #3]`” and “`STR r2, [r0, #3]`” are supported in Cortex-M, their access speed is slower than the aligned accesses “`LDR r1, [r0, #4]`” and “`STR r2, [r0, #4]`”.

10.2 Special Variables

This section discusses two special types of variables in C: static variables and volatile variables.

10.2.1 Static Variables

A static variable is initialized only once. Its lifetime is across its entire program runtime.

Different from local variables, a **static** C variable has a lifetime over the entire program runtime. A static variable declared within a C function is initialized only once at the compiling time no matter how many times this function is called. This static variable is visible only inside this function.

A static variable can be either global or local. A static local variable can only be visitable or available within the scope of the function in which this variable is declared. A static global variable can only be accessed within the source file in which it was declared, and other source files cannot access it.

A static variable is preferred to a global variable in C because a local or global static variable has a narrower access range. A program should avoid global variables whenever possible. A global variable is accessible to all source codes in any file of a software system. The major problem of using global variables is that they create hidden coupling between different software modules that is difficult to identify and understand, thus increasing the risk of software bugs. Because of the implicit interference created by global variables, a bug in one software module might cause the failure of another seemly-unrelated module, making the debug process difficult.

Always avoid global variables.

One effective way to avoid global variables is to use static variables. As shown in Example 10-5, the counter is declared as a global static variable, instead of a global variable. Therefore, the counter can only be accessed within in that source file. Codes in other source files cannot access the counter variable.

If a global variable is only accessed by one subroutine, the program may declare this variable as a local static variable within that subroutine. The access scope of a local static variable is the subroutine that declares it. When that subroutine is called successively, the value of the local static variable is retained.

```

static int counter = 0;

void increase(void){
    counter++;
}

void decrease(void){
    counter--;
}

```

Example 10-5. Example of using a global static variable

Example 10-6 and Example 10-7 compare how a local static variable and a local regular variable are accessed in assembly. All static variables are allocated in the data memory. However, a local variable is often stored in a register or the heap region of the data memory. A static variable is always loaded from the memory first, and then is stored back to the memory before exiting the subroutine. Therefore, if the subroutine is called again, the static variable keeps its previous value, instead of its initial value.

Example of a local non-static variable

C Program	Assembly Program
<pre> int foo(); int main(void) { int y; y = foo(); // y = 6 y = foo(); // y = 6 y = foo(); // y = 6 while(1); } int foo() { int x = 5; // x is a local variable x = x + 1; return(x) } </pre>	<pre> AREA static_demo, CODE EXPORT __main ALIGN ENTRY __main PROC BL foo ; r0 = 6 BL foo ; r0 = 6 BL foo ; r0 = 6 stop B stop ENDP foo PROC MOV r0, #5 ADD r0, r0, #1 BX lr ENDP END </pre>

Example 10-6. If x is not declared as static, foo() always returns the same value.

In the program given in Example 10-6, variable x is not declared as static. Thus, foo returns the same value each time it is called. From the assembly implementation, we can notice that the local variable x is always reinitialized when foo is called. Plus, in this example, x is stored in a register and its value is lost (not saved in the data memory) after foo exits.

Example of a local static variable

C Program	Assembly Program
<pre>int foo(); int main(void) { int y; y = foo(); // y = 6 y = foo(); // y = 7 y = foo(); // y = 8 while(1); } int foo() { // local static variable // x is initialized only once static int x = 5; x = x + 1; return(x) }</pre>	<pre>AREA myData, DATA ALIGN // Reserve space for x DCD 5 AREA static_demo, CODE EXPORT __main ALIGN ENTRY _main PROC BL foo ; r0 = 6 BL foo ; r0 = 7 BL foo ; r0 = 8 stop B stop ENDP foo PROC ; load address of x LDR r1, =x ; load value of x LDR r0, [r1] ADD r0, r0, #1 ; save value of x STR r0, [r1] BX lr ENDP END</pre>

Example 10-7. When *x* is declared as static, the *foo()* function returns different values.

In the program given in Example 10-7, variable *x* is declared as static locally within the *foo* function.

- The local static variable *x* is only initialized once. The initialization is carried out at compile time instead of at runtime. As you see from the assembly code, the variable is defined in the data region with an initial value of 5. No matter how many times *foo* runs, variable *x* is never re-initialized.
- When *foo* needs to increase the value of the local static variable *x*, the value of variable *x* is read from the memory at the beginning of *foo* and is saved into the memory before *foo* exits. Therefore, *foo* returns a different result each time it runs. On the contrary, *foo* in Example 10-6, in which *x* is not static, always returns the same value.

Example 10-8 gives another example of using static variables. The program uses the static variable *sum* to check whether an integer number is a palindrome number. A palindrome number remains the same if all digits are reversed.

C Program	Assembly Program
<pre> int isPal(int); int main(){ int n; n = isPal(12321); while(1); } // Check palindrome number int isPal(int n){ static int sum = 0; int r; if(n!=0) { r = n % 10; sum = sum*10 + r; isPal (n/10); } if (sum == n) return 1; else return 0; } </pre>	<pre> AREA myData, DATA ALIGN sum DCD 0 AREA palindrome, CODE EXPORT __main ALIGN ENTRY __main PROC LDR r0, =12321 BL isPal stop B stop ENDP ; Recursively check palindrome isPal PROC PUSH {r4, lr} MOV r4, r0 CBZ r4, done ; if n is 0, done MOV r2, #10 SDIV r1, r4, r2 ; r1 = n/10 MLS r3, r1, r2, r4 ; r3 = n - r1 * 10; LDR r1, =sum LDR r1, [r1] ; r1 = sum ADD r1, r1,r1,LSL #2 ; r1 = 5*sum ADD r1, r3,r1,LSL #1 ; sum = sum*10+r; LDR r2, =sum STR r1, [r2] ; save sum MOV r2, #10 SDIV r0, r4, r2 ; r0 = n/10 BL isPal ; recursive call done LDR r1, =sum LDR r1, [r1] CMP r1, r4 BNE no yes MOV r0, #1 ; if palindrome B exit no MOV r0, #0 ; if not palindrome exit POP {r4, pc} ENDP END </pre>

Example 10-8. Example of using a local static variable *sum* in a function

10.2.2 Volatile Variables

When the compiler optimizes a C program, a hard-to-find hidden error is that the program mistakenly reuses the value of a variable stored in a register, instead of reloading it from memory each time. In order to avoid such compilation error, the program should declare the variable as volatile, such as:

```
volatile int variable;
```

The keyword volatile forces the compiler to generate an executable, which always loads the variable value from the memory whenever this variable is read, and always stores the variable in the memory whenever it is written.

A variable declared as **volatile** indicates that this variable may be changed by an input, not by the program.

Therefore, the program should always re-read its value from the memory if the variable is used.

Example 10-10 gives a simple example to illustrate the necessity of declaring a variable counter, shared by two concurrently running tasks (*main* function and *SysTick_Handler*), as volatile. In this example, the main program uses the SysTick to implement a time delay. It sets up the SysTick timer and then waits until the SysTick interrupt service routine reduces the counter to 0. The SysTick decrements the counter by one when a system timer interrupt occurs. Chapter 11.4 gives the implementation of *SysTick_Init()*.

Main Program (<i>main.c</i>)	Interrupt Service Routine (<i>isr.s</i>)
<pre>// volatile unsigned int counter; unsigned int counter; extern void task(); extern void SysTick_Init(); int main(void) { counter = 10; SysTick_Init(); while(counter != 0); // Delay // Continue the task while(1); }</pre>	<pre>AREA ISR, CODE, READONLY IMPORT counter ENTRY SysTick_Handler PROC EXPORT SysTick_Handler LDR r1, =counter LDR r0, [r1] ; load counter SUB r0, r0, #1 ; counter-- STR r0, [r1] ; save counter BX LR ; exit ENDP END</pre>

Example 10-9. A C variable is not declared as *volatile* while it should be.

Compilers often attempt to optimize the program, but sometimes can cause troubles. The compiler observes that, after the counter is initialized to 10, the value of the counter variable is not modified directly by *main()* or indirectly by any subroutine called from *main()*. Loading data from memory is much slower than retrieving data from registers.

Therefore, the compiler may choose to reuse the value of the counter stored in a register, instead of fetching the counter value again from the memory, when the counter is accessed in the while loop. Example 10-10 compares the assembly program generated by the compiler when the counter variable is declared as volatile or non-volatile.

- If the counter is not declared as volatile, the while loop is a dead loop. *SysTick_Handler* periodically decrements the counter and stores its value in the memory. However, the main program repeatedly checks register r0, without reloading the latest value of the counter from the memory.
- If the counter is declared as volatile, the dead loop problem is avoided.

If counter is not declared as volatile		If counter is declared as volatile	
main PROC		main PROC	
LDR r1, =counter		LDR r1, =counter	
MOV r0, #10		MOV r0, #10	
STR r0, [r1]		STR r0, [r1]	
BL SysTick_Init		BL SysTick_Init	
wait	CMP r0, #0 ; r0 does not hold ; latest counter value	wait	LDR r1, =counter LDR r0, [r1]
	BNE wait ; Thus, a dead loop		CMP r0, #0
stop	B stop	stop	BNE wait
	ENDP		B stop
			ENDP

Example 10-10. Comparison of assembly instructions generated by the compiler when the counter variable is declared as volatile and non-volatile.

A C program should declare any variable that represents the data of a memory-mapped I/O register as **volatile**. Memory-mapped I/O has been widely used to access peripheral devices. Data and control registers of external devices are mapped to specific memory addresses, and a program can use memory pointers to access these hardware registers, such as the following C statement.

```
unsigned int *p = (unsigned int *) 0x60002400;
```

To prevent the compiler from optimizing out these memory pointers incorrectly, these pointers must be declared as volatile. The following example uses a memory pointer to access a 32-bit hardware register mapped to the memory address 0x60002400.

```
volatile unsigned int *p = (unsigned int *) 0x60002400;
```

In sum, a variable should be declared as volatile to prevent the compiler from optimizing it away when (1) this variable is updated by external memory-mapped hardware, or (2) this variable is global and is changed by interrupt handlers or by multiple threads.

10.3 Inline Assembly

A block of assembly code, called inline assembly, can be directly embedded in a C program. It is convenient for programmers because it does not require different assemble and link processes. Another advantage of inline assembly is that it can flexibly access C variables without export and import operations, which would be needed if the assembly code were written as an assembly subroutine.

10.3.1 Assembly Functions in a C Program

When a C program declares a function with “`_asm`”, the assembly code in this function has to preserve the environment.

When a block of assembly code is embedded within a function by using “`_asm`”, the assembly code does not need to preserve the environment.

A C program can have inline assembly by using the “`_asm`” keyword. It has two different uses. The first is to specify a function that is implemented in assembly completely. The second is to specify multiple lines of assembly code within a C function.

When a function is declared with “`_asm`”, the assembly implementation must preserve the runtime environment via the stack and recover the environment before exiting from the subroutine. The assembly code can directly access the registers and must follow the procedure call protocol. Example 10-11 and Example 10-12 use “`_asm`” to implement a C function in assembly.

```
_asm int sum4(int a, int b, int c, int d){
    ; arguments stored in r0, r1, r2, r3
    PUSH {r4, lr}          ; preserve environment in stack
    MOV r4, r0              ; r0 = 1st argument
    ADD r4, r4, r1           ; r1 = 2nd argument
    ADD r4, r4, r2           ; r2 = 3rd argument
    ADD r0, r4, r3           ; r3 = 4th argument, r0 = return
    POP {r4, pc}             ; recover environment from stack
}

int main(void){
    int s = sum4(1, 2, 3, 4);
    while(1);
}
```

Example 10-11. Using inline assembly to implement a subroutine that adds four integers.

```

char a[25] = "Hello!";
char b[25];

__asm void strcpy(char *src, char *dst){
loop LDRB r2, [r0], #1 ; 1st argument, r0 = src, post-index
    STRB r2, [r1], #1 ; 2nd argument, r1 = dst, post-index
    CMP r2, #0
    BNE loop
    BX lr
}

int main(void){
    strcpy(a, b);
    while(1);
}

```

Example 10-12. Using inline assembly to copy a string.

When a function is declared with “`__asm`”, the compiler only creates the interface of this function and does not provide any actual implementation. Therefore, Example 10-11 uses `PUSH` and `POP` to preserve and recover the running environment of the caller.

10.3.2 Inline Assembly Instructions in a C Program

When “`__asm`” is used to declare a block of assembly instructions in a C function, the assembly code cannot access registers and does not need to preserve the runtime environment in the stack. The compiler automatically generates necessary code to preserve the environment.

Additionally, the assembly code treats each C variable as a register. These C variables are called *virtual registers* and they can be accessed in assembly instructions. Compilers replace virtual registers with real registers.

Also, the comments of the assembly code have to be in C style. Example 10-13 has a block of assembly instructions in a C function.

```

int sum4(int a, int b, int c, int d){
    int t;
    __asm {
        ADD t, a, b; // t, a, and b are virtual registers
        ADD t, c;    // Cannot directly access r0 - r15
        ADD t, d;    // Have to use comment style of C
    }
    return t;
}
int main(void){
    int s = sum4(1, 2, 3, 4);
    while(1);
}

```

Example 10-13. Using “`__asm`” to declare a block of assembly instructions in a C function.

10.4 Calling Assembly Subroutines from a C Program

A large software project often has its program code saved in multiple small source files, instead of a single monolithic file. This technique not only improves the software modularity and maintainability, but also reduces the compilation time. These files can be compiled separately so that unmodified files do not need to be recompiled.

This section shows how a C program calls assembly subroutines that are in separate source files.

- The assembly code must use the directive “EXPORT” or “GLOBAL” to make all variables or subroutines that are accessed in the C program as global. These directives make subroutine names visible outside this source code module. Consequently, the compiler can locate them when linking the object files generated from their source codes.
- The C program must declare these functions by using the keyword “extern”.

10.4.1 Example of Calling an Assembly Subroutine

In the following example, the C program calls the assembly subroutine *strlen*, which calculates the length of a string. The C program and the assembly program are in two separate source files: *main.c* and *strlen.s*.

C Program (<i>main.c</i>)	Assembly Program (<i>strlen.s</i>)
<pre>char str[25] = "Hello!"; extern int strlen(char* s); int main(void){ int i; i = strlen(str); while(1); }</pre>	<pre>AREA stringLength, CODE EXPORT strlen ; make strlen visible ALIGN strlen PROC PUSH {r4, lr} ; preserve r4 and lr MOV r4, #0 ; initialize length loop LDRB r1, [r0, r4] ; r0 = string address CBZ r1, exit ; branch if zero ADD r4, r4, #1 ; length++ B loop ; do it again exit MOV r0, r4 ; place result in r0 POP {r4, pc} ; exit ENDP</pre>

Example 10-14. A C program calls an assembly routine stored in a different file.

- The assembly subroutine follows the procedure call protocol defined in ARM embedded application binary interface (EABI) and assumes argument *str* is passed in register *r0*. Furthermore, the caller expects that the assembly subroutine returns a 32-bit result in register *r0* and a 64-bit result in registers *r1:r0*.

- The C program declares the assembly function to be called by using the keyword "extern" to inform the compiler that the implementation of this function is in another file.
- The assembly subroutine uses "EXPORT strlen" to make the symbol strlen visible to the linker. Note all symbols are case-sensitive.

10.4.2 Example of Accessing C Variables in Assembly

An assembly program can access global variables defined in a C program or a separate assembly source file. When an assembly program accesses a global variable defined elsewhere, it needs to import that variable name by using the directive "IMPORT". An imported variable name, or called a symbol, is resolved at link time. In the following example, the C program declares the global variable counter. The assembly code uses "IMPORT counter" to access this global variable.

C Program (main.c)	Assembly Program (count.s)
<pre>int counter; extern int getValue(); extern void setValue(int c); void increment(); int main(void) { int c = 0; setValue(1); increment(); c = getValue(); while(1); } void increment(){ counter += 1; }</pre>	<pre>AREA count, CODE IMPORT counter ALIGN setValue PROC EXPORT setValue LDR r1, =counter STR r0, [r1] BX lr ENDP getValue PROC EXPORT getValue LDR r1, =counter LDR r0, [r1] BX lr ENDP increment PROC EXPORT increment [WEAK] LDR r1, =counter LDR r0, [r1] ADD r0, r0, #1 STR r0, [r1] BX lr ENDP END</pre>

Example 10-15. Example of accessing a C variable in assembly routines

The assembly program exports the *increment* symbol with weak specified. By default, all exporting statements are strong. At the linking stage, a strongly exported symbol replaces

a weakly exported symbol of the same name. The linker reports a fatal error if there are more than one strong instance of the same symbol name. Because the symbol *increment* is exported weakly and the one in the C program is exported strongly, the *increment* function defined in C overrides the one defined in the assembly. Consequently, when the *increment* function is called, the *counter* variable is incremented by two, instead of one.

10.5 Calling C Functions from Assembly Programs

An assembly program can call functions implemented in C. The assembly program needs to follow the procedure call protocol defined in ARM embedded application binary interface (EABI).

- Specifically, the assembly program needs to place the input arguments of a C function in registers r0-r1 before it calls the function.
- The assembly program also expects that the result is returned in register r0 if the C function returns a value less than 32 bits.
- If the result has more than 32 bits, registers r0 – r4 can be used to hold the result.

10.5.1 Example of Calling a C Function

In the following example, the assembly program calls the *strlen* function implemented in C. The C function returns the length of the string in register r0 to the assembly program. In the assembly code, the C function names have to be imported to avoid linking errors.

Assembly Program (<i>main.s</i>)	C Program (<i>strlen.c</i>)
<pre> AREA my_strlen, CODE EXPORT __main IMPORT strlen ALIGN ENTRY __main PROC LDR r0, =str BL strlen stop B stop ENDP AREA myData, DATA ALIGN Str DCB "12345678",0 END </pre>	<pre> int strlen(char *s){ int i = 0; while(s[i] != '\0') i++; return i; } </pre>

Example 10-16. Example of an assembly program that calls a C subroutine

If "WEAK" is specified in the import directive, the linker does not produce any error if the symbol is not defined externally. Instead, the linker then replaces the symbol with zero or some appropriate value. For example, if the label is not defined in the project, the linker then replaces it with the address of the next instruction after the branch.

```
IMPORT label [WEAK]
...
B label
...
```

Example 10-17. A symbol declared with weak prevents the linker from fatal linking error.

10.5.2 Example of Accessing Assembly Data in a C Program

To access variable *counter* defined in the assembly code, the C program has to use "extern int counter". This statement is to inform the compiler that this variable is defined outside this C program.

The assembly program allocates the memory space for the counter variable, and the C program only needs to indicate the existence of this variable without performing any memory allocation. Without the "extern" keyword, compilers would allocate space again for variable *counter* defined in the C program, thus producing an error of duplicated variables at the link stage.

Assembly Program	C Program
<pre>AREA main, CODE EXPORT __main IMPORT getValue IMPORT increment IMPORT setValue ALIGN ENTRY _main MOVS r2, #0 MOVS r0, #1 BL setValue BL increment BL getValue MOV r2, r0 stop B stop AREA myData, DATA EXPORT counter counter DCD 0 END</pre>	<pre>extern int counter; int getValue() { return counter; } void increment() { counter++; } void setValue(int c) { counter = c; }</pre>

Example 10-18. C functions access a variable defined in an assembly program.

10.6 Exercises

1. Translate the following C statement into assembly.

```
x[1].c += 100;
```

The following gives the definition the structure array $x[2]$. Assume register r_0 holds the starting address of this array.

```
__packed struct X {
    uint8_t a;
    int32_t b;
    int16_t c;
    int32_t d;
    uint8_t e;
} x[2];
```

2. Translate the following C statement into assembly.

```
x[1].c += 100;
```

The structure array is defined below. Assume register r_0 contains the starting address of this array.

```
struct Y {
    uint8_t a;
    int32_t b;
    int16_t c;
    int32_t d;
    uint8_t e;
} y[2];
```

3. Write a subroutine in assembly that removes all occurrences of a given character in a string. The subroutine takes two parameters: the string pointer, and the character to be deleted. Write a C code that calls this subroutine. The string is defined as global in the C code.
4. Suppose we have the following *strcat* function written in C, which concatenates the second string to the first string. Write an assembly program that calls the *strcat* function. These two strings are defined in the data area in the assembly code.

```
void strcat (char * dst, char * src) {
    while(*dst++);
    while(*dst++ = *src++);
}
```

5. Write a subroutine *swap* in assembly that swaps two strings, and a C program that calls the swap subroutine. These two strings are defined in the C program. (Hint: There is no need to swap all characters in the strings, and swapping the memory pointers in the assembly is sufficient.)
6. Write an assembly program that calls the following C subroutine that returns the memory address of the last occurrence of a given character in a string.

```
char * search (char * s, char c) {
    char *p = NULL;
    for(; *s; s++)
        if (*s == c)
            p = s;
    return p;
}
```

7. Suppose the following structure array is defined as global in a C program. Write an assembly program that iterates through the array and find the total scores.

```
struct Student_T {
    char c1;
    char c2;
    int score;
    char c3;
} students[10];
```

8. Write a subroutine *max4* in assembly to find the maximum value among four signed integers. These integers are passed to the subroutine via registers. Write a C program to test the *max4* subroutine.
9. Write an assembly subroutine that checks whether a given integer is a palindrome number. For example, 9, 11, 1234321, 141, 1221, and 120021 are palindrome numbers. Write a C program that calls the assembly subroutine. The input is an unsigned integer. The return is 1 if the number is a palindrome and 0 if not.
10. Identify and correct the errors in the following inline assembly program that calculates the sum of four integers.

```
_asm int sum4(int a, int b, int c, int d)
// arguments stored in r0 - r3
MOV r4, r0      ; r0 = 1st argument
ADD r4, r4, r1  ; r1 = 2nd argument
ADD r4, r4, r2  ; r2 = 3rd argument
ADD r0, r4, r3  ; r3 = 4th argument, r0 = return
}
```

```
int main(void){
    int s = sum4(1, 2, 3, 4);
    while(1);
}
```

11. Identify and correct the errors in the following inline assembly codes that calculate the sum of four integers.

```
int sum4(int a, int b, int c, int d){
    int t;
    __asm {
        ADD t, r0, r1;
        ADD t, r2;
        ADD t, r3;
    }
    return t;
}

int main(void){
    int s = sum4(1, 2, 3, 4);
    while(1);
}
```

12. Identify and correct the errors in the following two programs.

C Program (main.c)	Assembly Program (strcpy.s)
<pre>char src[9] = "Hello!"; char dst[9]; int main(void){ strcpy(dst, src); while(1); }</pre>	<pre>AREA stringCopy, CODE ALIGN strcpy PROC loop LDRB r2, [r1] ; Load a byte, r1 = *src STRB r2, [r0] ; Store a byte, r0 = *dst ADD r1, #1 ; Increase memory pointer ADD r0, #1 ; Increase memory pointer CMP r2, #0 ; Zero terminator BNE loop ; Loop if not null terminator ENDP END</pre>